

## GPU Computing is everywhere!



[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

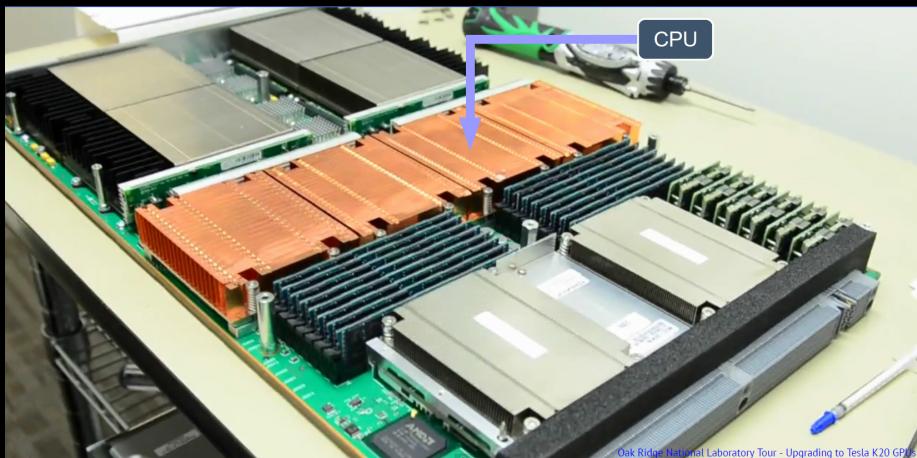
## GPU in a HPC system



Oak Ridge National Laboratory Tour - Upgrading to Tesla K20 GPUs

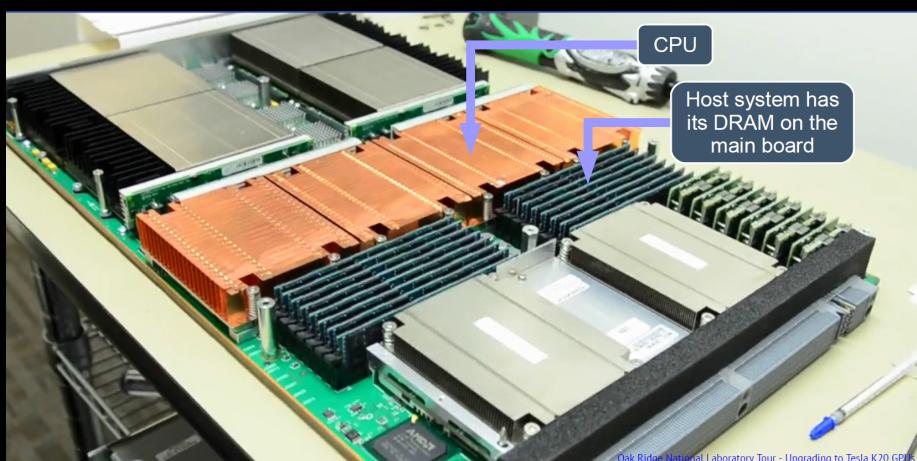
[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

## GPU in a HPC system



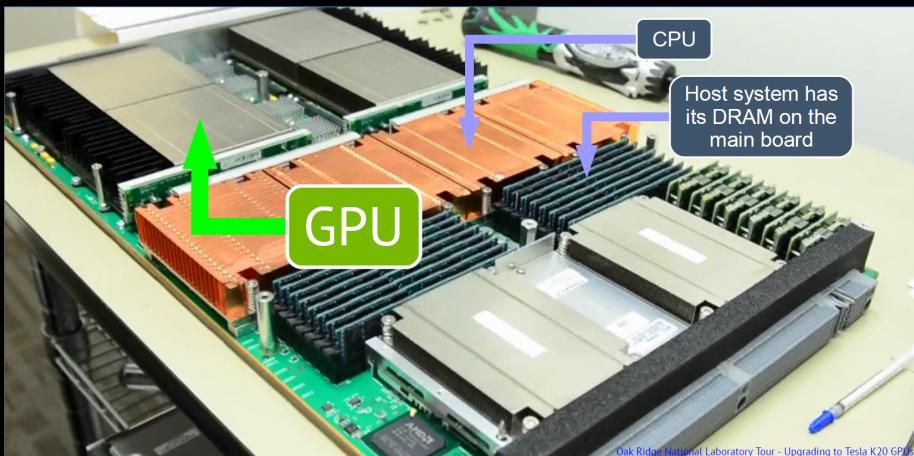
[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

## GPU in a HPC system



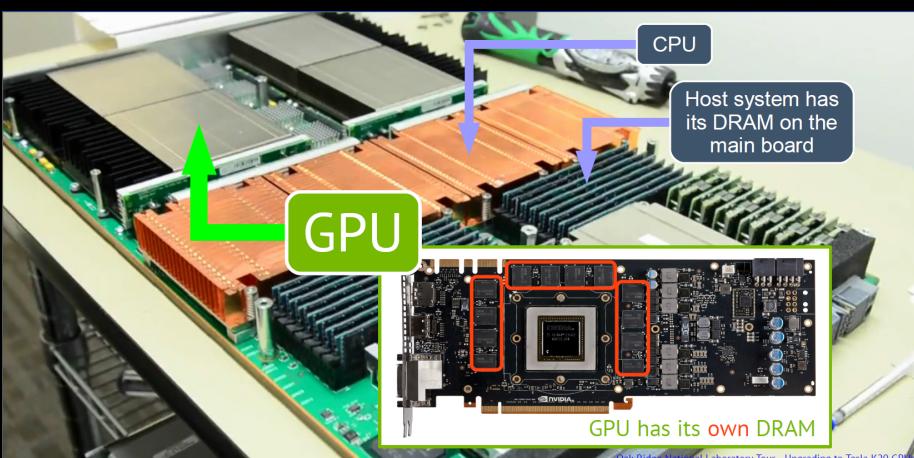
[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

## GPU in a HPC system



[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

## GPU in a HPC system

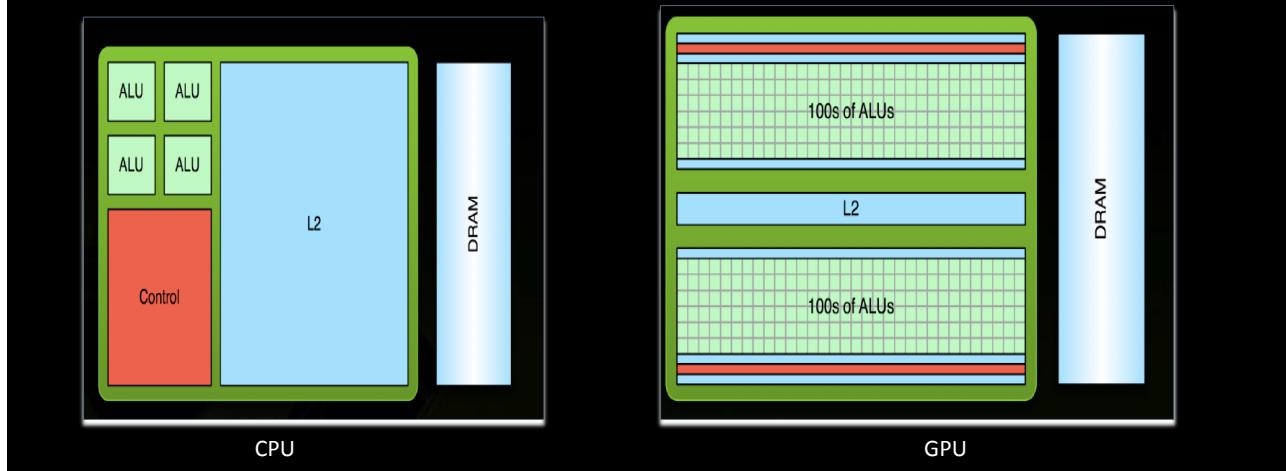


[http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day\\_1](http://www.hpcadvisorycouncil.com/events/2014/swiss-workshop/presos/Day_1)

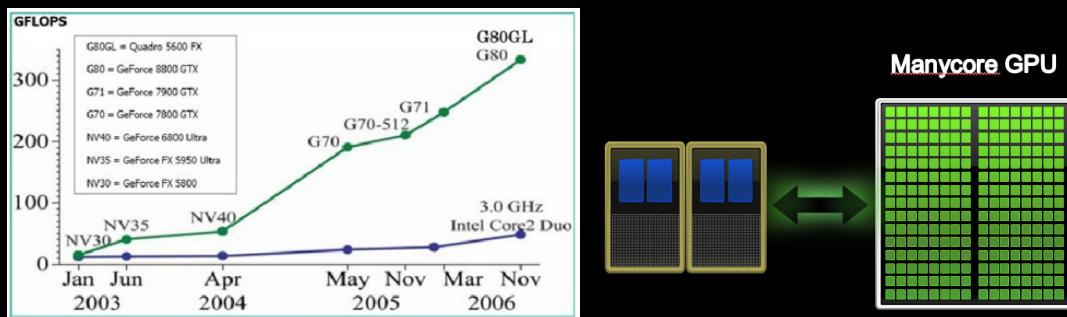
Why do we need a GPU????

## Difference between a CPU and GPU

- Both cater to different needs → **Low Latency or High Throughput?**



## Difference between Multicore CPU and Manycore GPU



<http://image.slidesharecdn.com/2009-07-22cuda-techtalk-090723122825-phpapp01/95/tech-talk-nvidia-cuda-7-728.jpg>

## HPC Machines in CS Labs

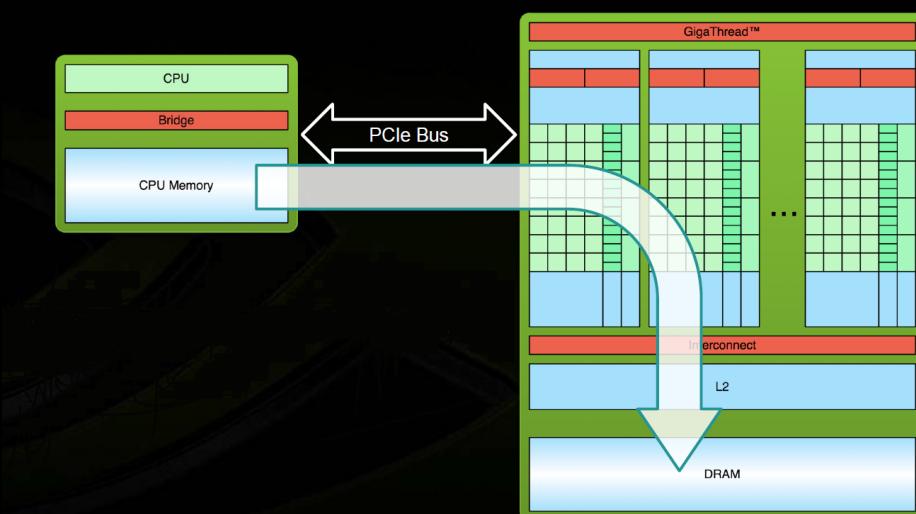
Name	GPUs
bentley	<b>GF119 [NVS 315] and GM204 [GeForce GTX 980]</b>
bugatti	<b>GF119 [NVS 315] and GM204 [GeForce GTX 980]</b>
ferrari	<b>GF119 [NVS 315] and GM204 [GeForce GTX 980]</b>
jaguar	<b>GF119 [NVS 315] and GM204 [GeForce GTX 980]</b>
lamborghini	<b>GF119 [NVS 315] and GM200 [GeForce GTX TITAN X]</b>
lotus	<b>GF119 [NVS 315] and GM200 [GeForce GTX TITAN X]</b>
maserati	<b>GF119 [NVS 315] and GM200 [GeForce GTX TITAN X]</b>
porsche	<b>GF119 [NVS 315] and GM200 [GeForce GTX TITAN X]</b>
raspberries	<b>G96GL [Quadro FX 580]</b>

## Features of the Graphics card **GeForce GTX 980**

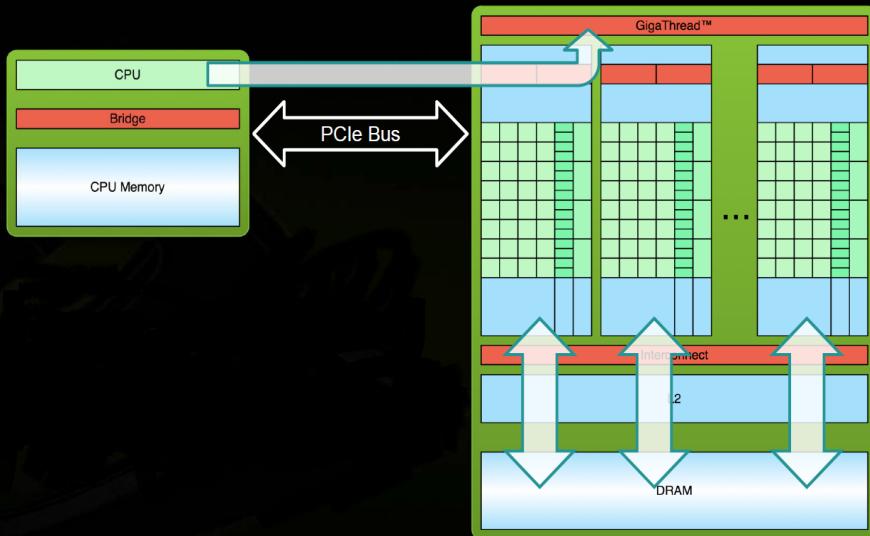
GPU Architecture	Maxwel
GPU Name	GM204
CUDA Cores	2048
Clock Speed	1126 MHz
VRAM	4 GB GDDR5
Memory Bus	256-bit
Memory Clock	7.0 GHz
Memory Bandwidth	224.0 GB/s
Power Connectors	Two 6-Pin



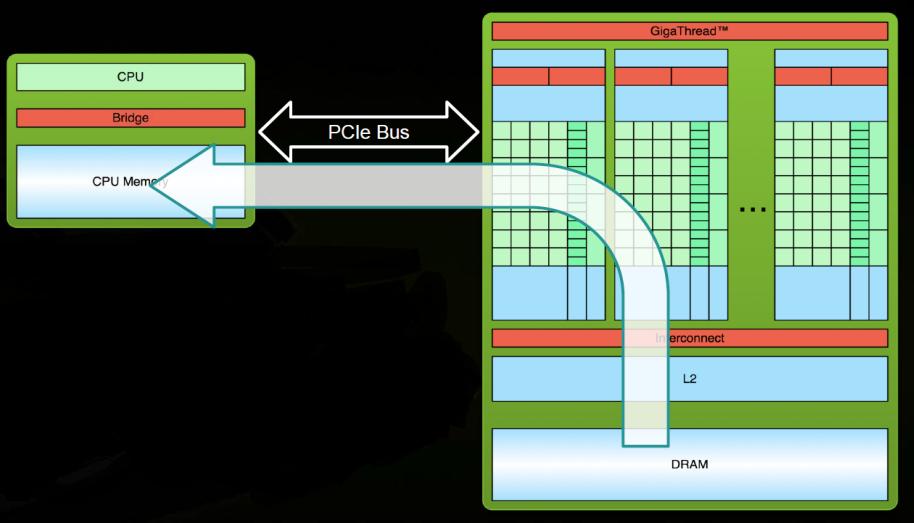
## Processing Flow:



## Processing Flow(contd.):



## Processing Flow(contd.):



# GPU Architecture

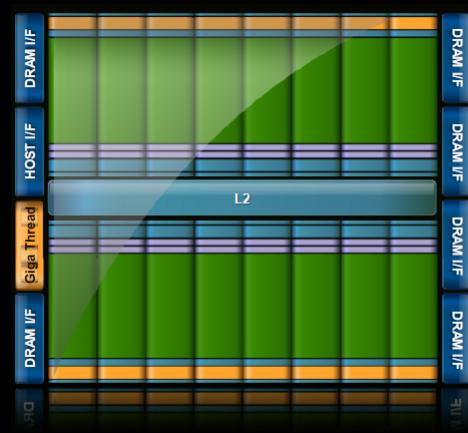
## GPU Architecture : Two Main Components

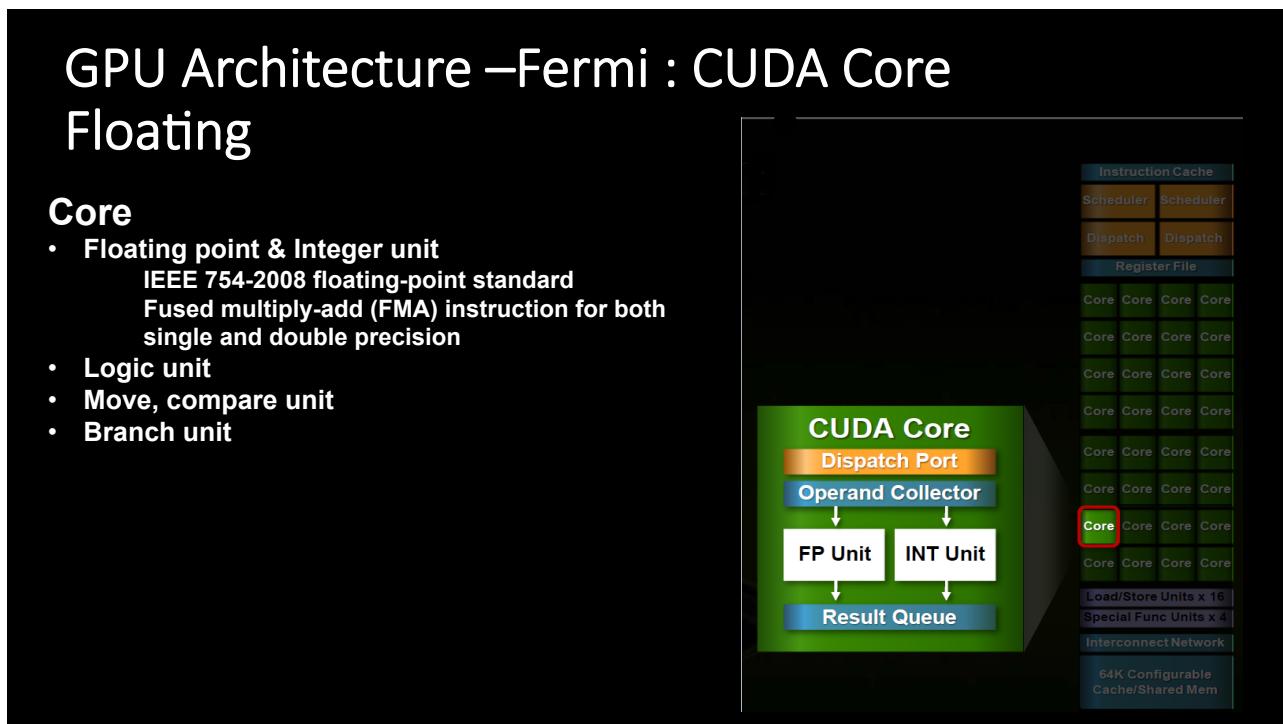
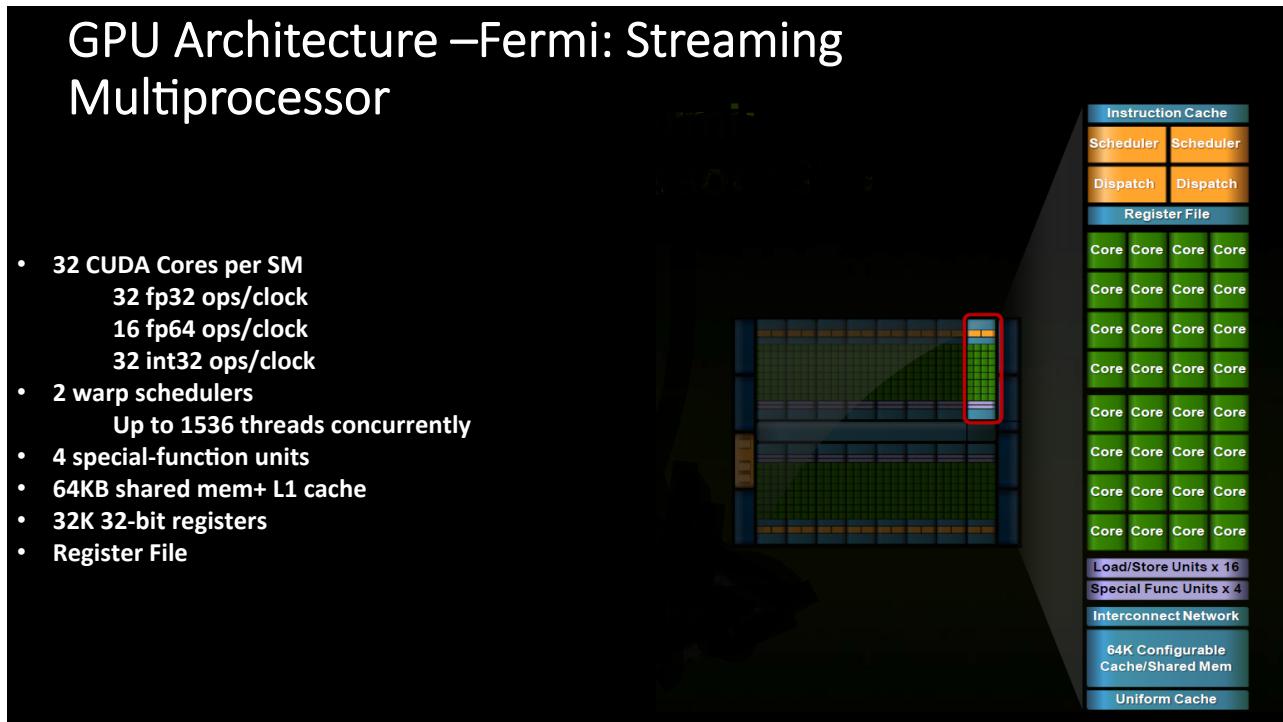
### 1. Global Memory:

- Accessible by both GPU and CPU
- Analogous to RAM in a CPU server

### 2. Streaming Multiprocessors (SMs):

- Perform the actual computations
- Each SM has its own:  
Control units, registers, execution pipelines, caches





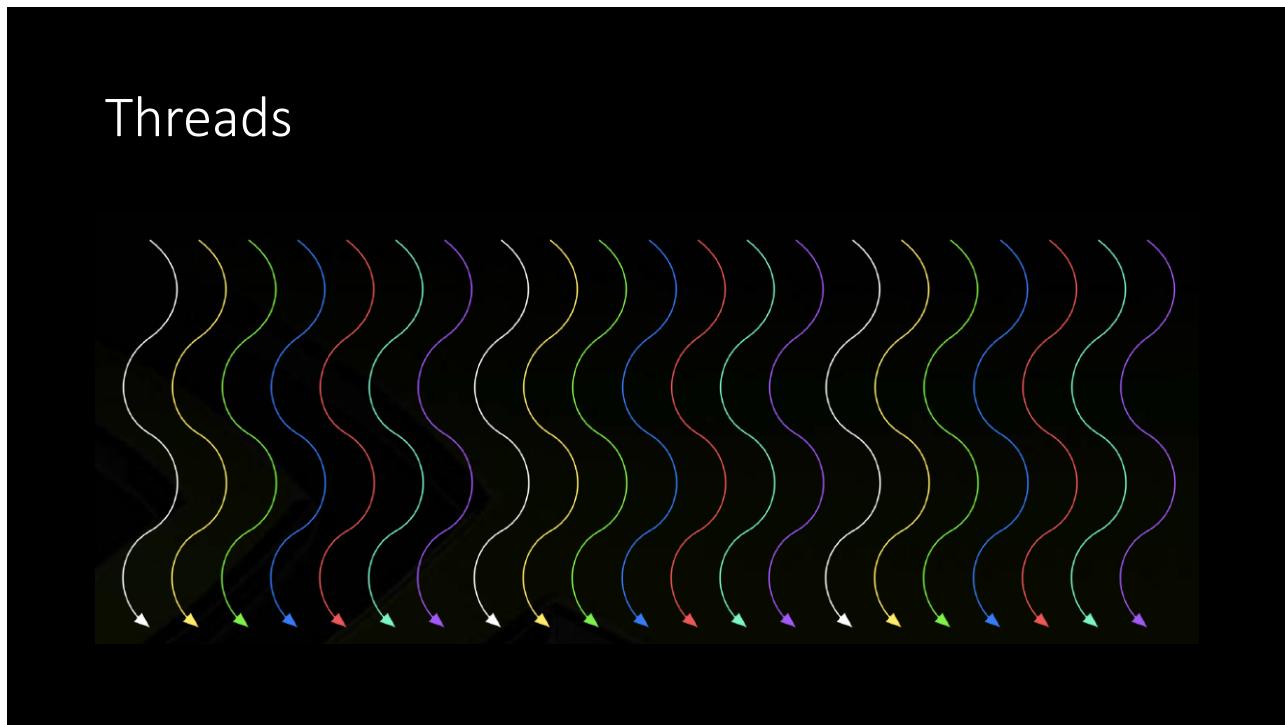
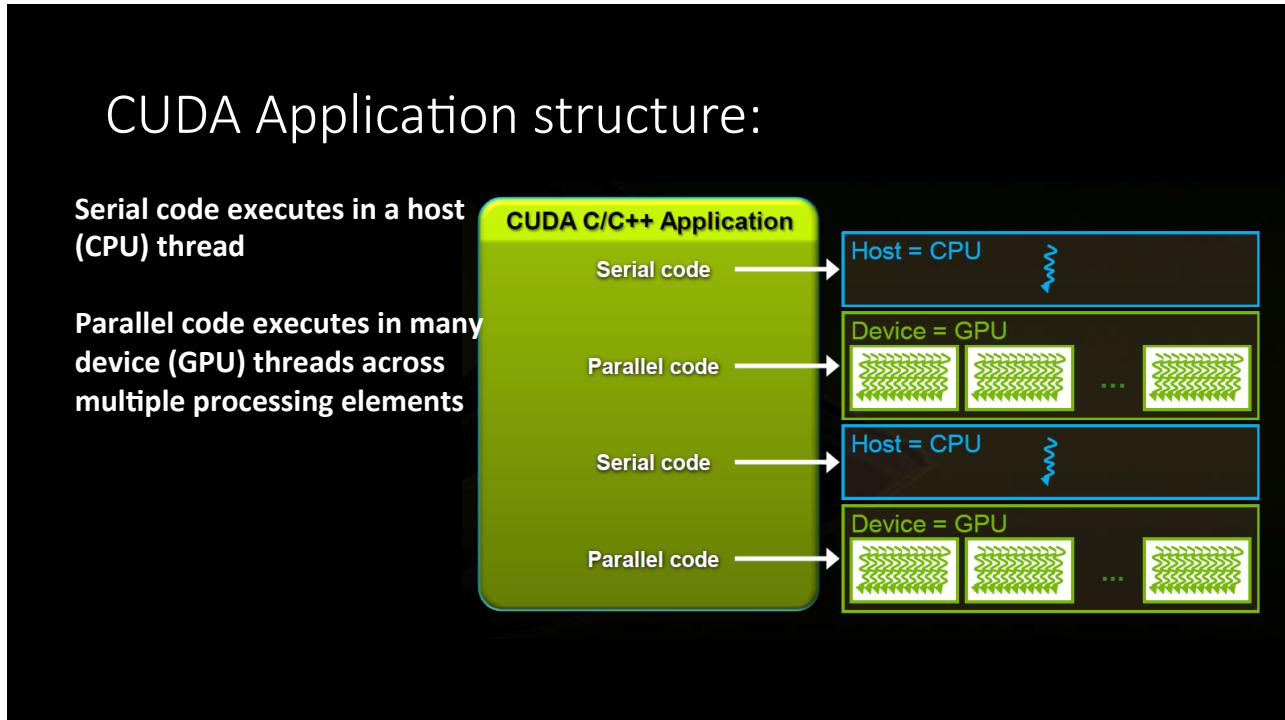
## Memory System-- Architecture

### **Shared memory(L1)**

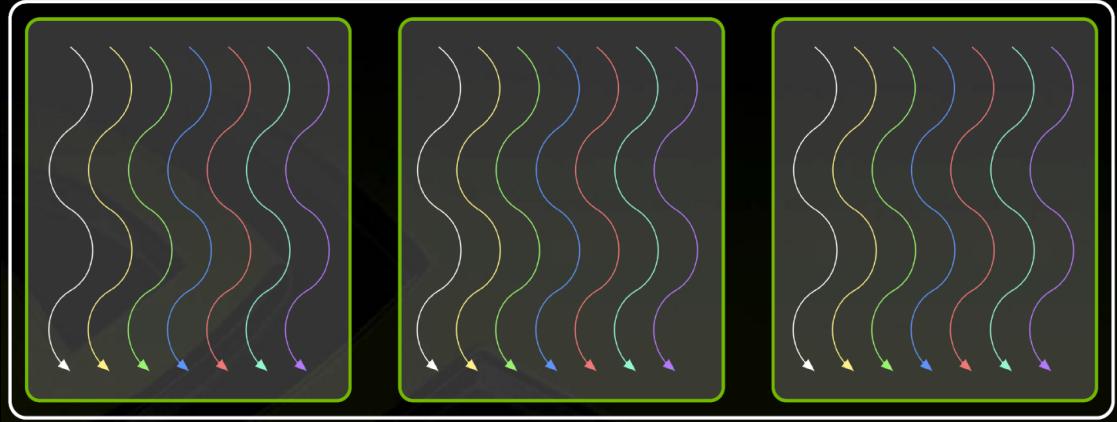
- User-managed scratch-pad
- Hardware will not evict until threads overwrite
- 16 or 48KB / SM (64KB total is split between Shared and L1)
- Aggregate bandwidth per GPU: 1.03 TB/s

ECC Protection can be enabled.

CUDA Programming  
Abstractions  
Shivani Dave

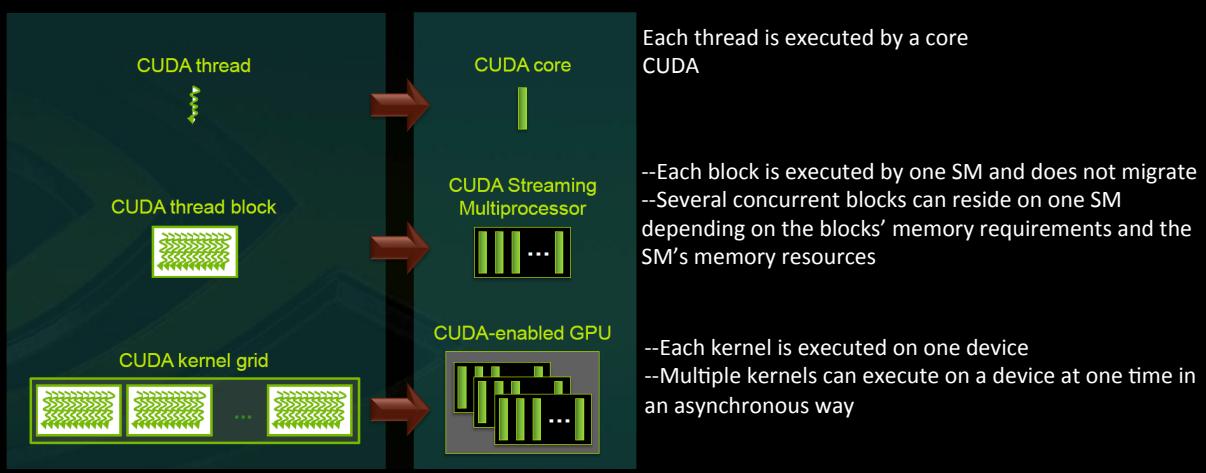


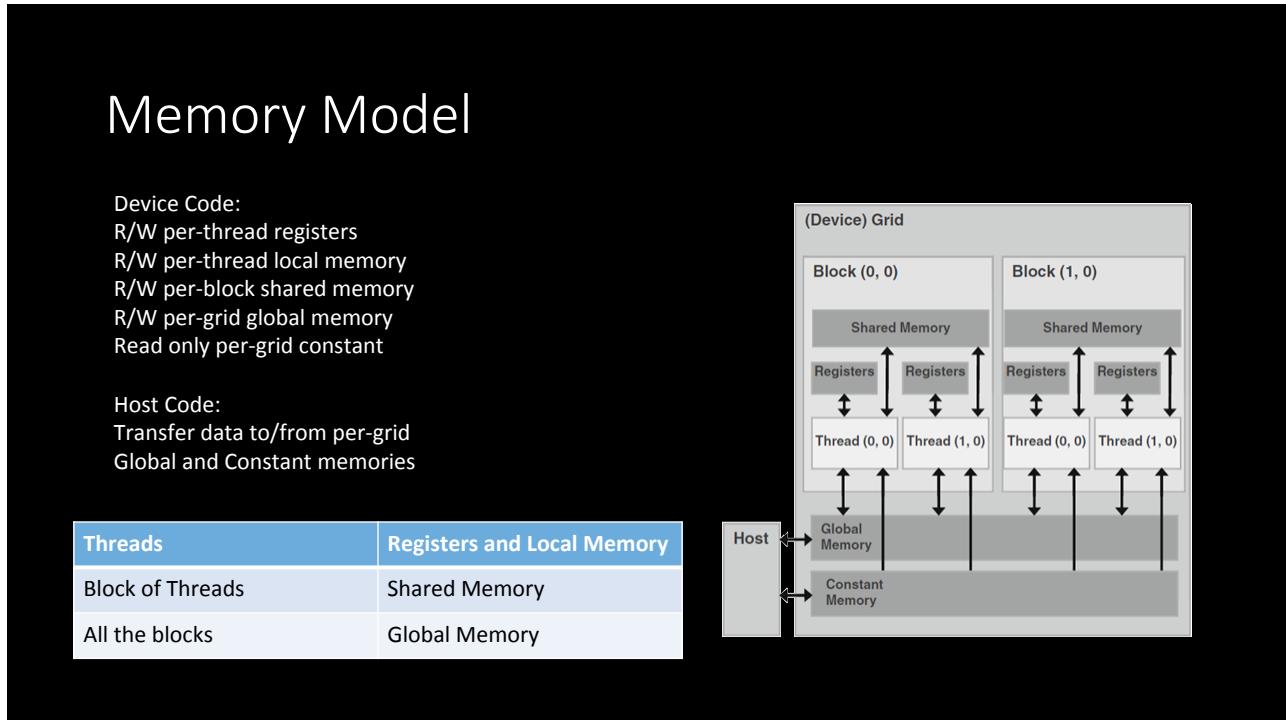
## Threads are grouped into blocks



- A kernel is executed as a grid of blocks of threads

## Kernel Execution





## Memory Model (contd ...)

- Global memory is the slowest memory on the GPU
- Coalescing improves memory performance; it occurs when multiple (row major order) consecutive threads (IDs) read / write consecutive data items from / to global memory
- 16 (half a warp) global array elements are accessed at once: coalescing produces vectorized reads / writes that are much faster than element wise reads / writes.
- A warp in CUDA, is a group of 32 threads, which is the minimum size of the data processed in SIMD fashion by a CUDA multiprocessor.

## Sequential C code -> Vector Addition

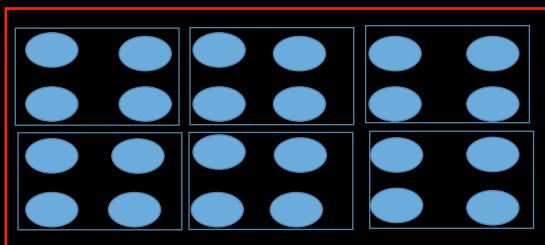
```
#include<stdio.h>
void Addition_Vector(int *a, int *b, int *c)
{
    int i=0;
    for(i=0;i<25;i++)
        c[i]=a[i]+b[i];
}

int main()
{
    int a[25],b[25],c[25],i;
    for(i=0;i<25;i++)
    {
        a[i]=i;
        b[i]=i+1;
    }

    Addition_Vector(a,b,c);
    for(i=0;i<25;i++)
        printf("\n%d",c[i]);
}
```

## Parallelism and Threads

- To process operations in parallel, the operations must be independent of each other i.e., no data dependencies
- A thread is mapped to a single processor which executes in parallel with the remaining threads.
- A CUDA kernel is executed by a grid (array) of threads with each thread with a unique index id in a specific block.
- A grid is organized as a 2D array of blocks (gridDim.x and gridDim.y)
- Each block is organized as 3D array of threads (blockDim.x, blockDim.y, and blockDim.z)



gridDim = (3,2)  
blockDim = (2,2,1)

## Thread Allocation

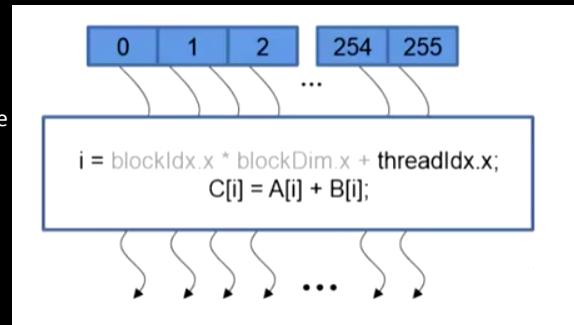
- A thread block can be allocated on any stream multiprocessor and thread blocks must be independent of each other, i.e., cannot communicate with each other at all.
  - pro: now the computation can run on any number of SMs
  - con: this makes programming a GPU harder
- Multiple thread blocks can be scheduled on one multiprocessor, if resources allow it. They still are independent of each other.

## Thread Synchronization

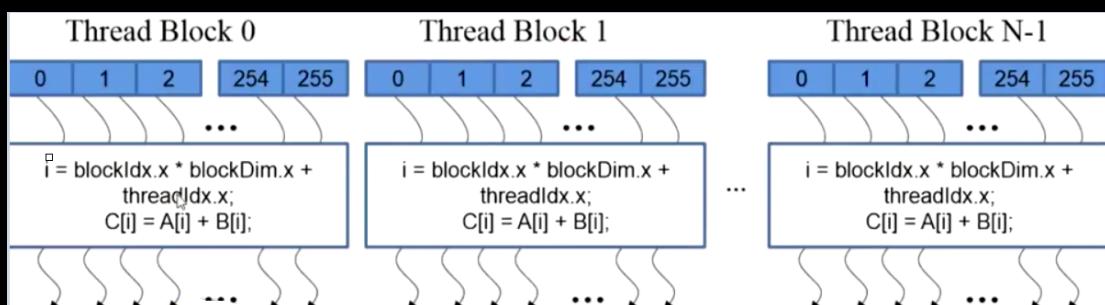
- Threads inside one thread block can synchronize – `_syncthreads()` command
- host can synchronize kernel calls – either explicitly through `cudaThreadSynchronize()` – or implicitly through `memcpy()`-s

## Parameterizing the thread code

- The programmer's job is to write the code so that when a collection of threadblocks, each with a collection of threads executes in concert, the collective work done solves the problem.
- But there is only one piece of code (what it does is a function of the grid and thread "coordinates."
  - GridIdx.x, GridIdx.y
  - threadIdx.x, threadIdx.y, threadIdx.z
- In our first example, there is just one block which has 256 threads in it: all but threadIdx.x are zero, and unused.
- So, we will have all values from 0 to 255
- Analogy with OpenMP/MPI:
  - Dim is line num\_threads(), and MPI\_Size
  - Ids is like thread\_num() and MPI\_rank



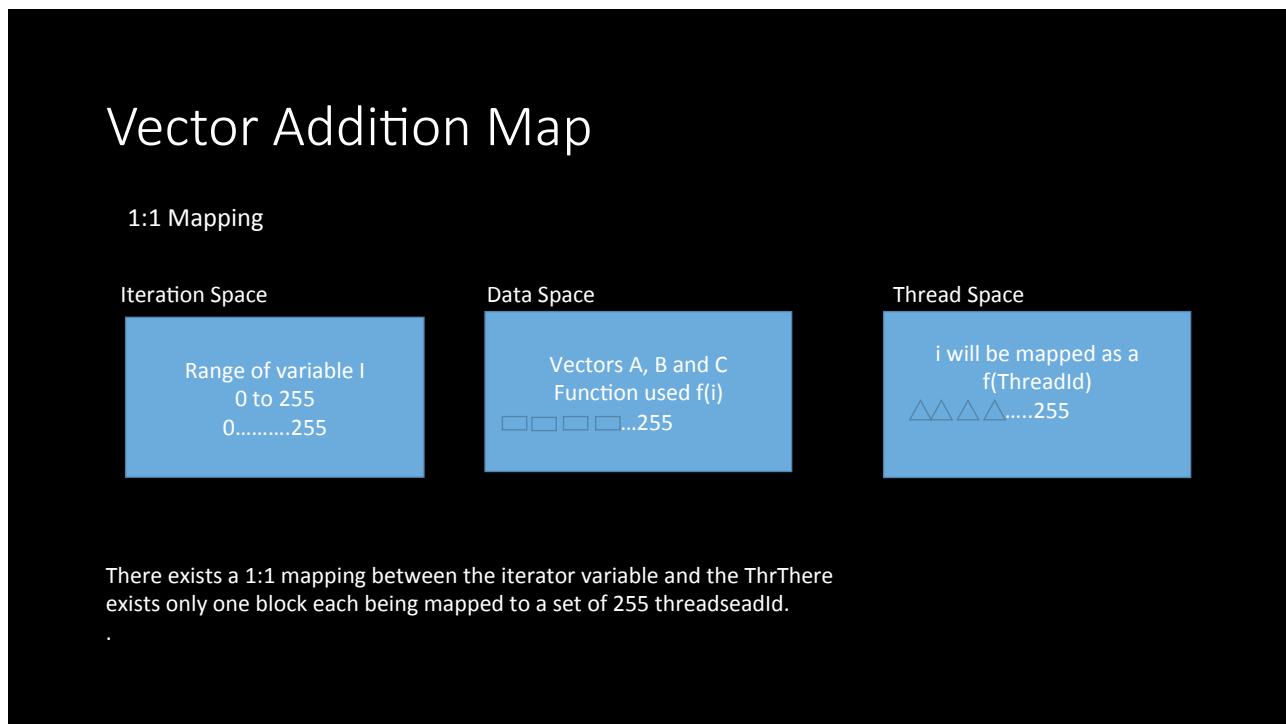
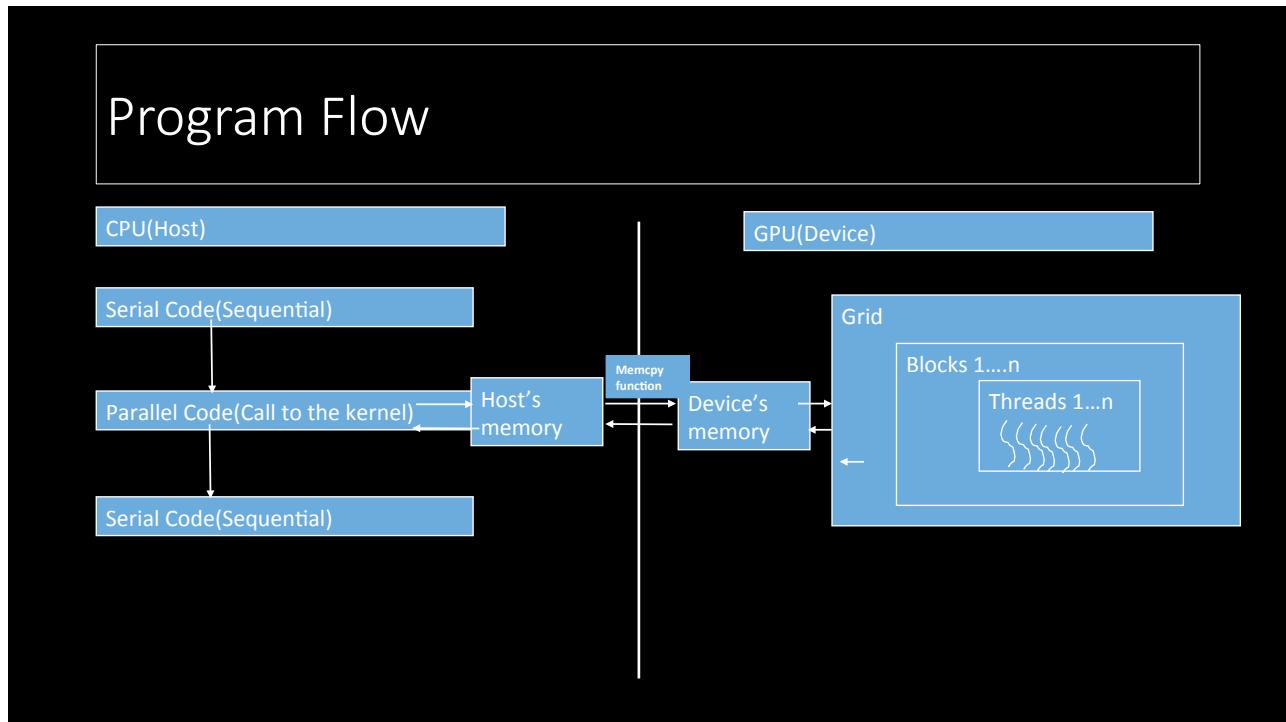
## Multiple Blocks



Multiple block in the Grid.

Each block will have an equal number of threads as defined in the program.

In the above example: blockIdx.x keeps on increasing from 0 to 255 similar to the threadIdx.x . The variable blockDim.x will be 256 which is equal to the number of blocks.



## Parallel Code → Vector Addition

```
#include<cuda.h>
void main(){
    int n=256; int b[n]; int a[n],c[n]; //initialize the arrays too
    int *dev_b; int *dev_a, int *dev_c;
    cudaMalloc((void**)&dev_a, n*sizeof(int));
    cudaMalloc((void**)&dev_b, n*sizeof(int));
    cudaMalloc((void**)&dev_c, n*sizeof(int));
    cudaMemcpy(dev_a, &a, n*sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, &b, n*sizeof(int), cudaMemcpyHostToDevice);
    dim3 gridDim = 1;//Number of blocks in the grid
    dim3 blockDim =n;//Number of threads in each block
    Addition_Vector<<<gridDim,blockDim>>>(dev_a,dev_b,dev_c); // Call to the device
    cudaMemcpy(&c, dev_c, n*sizeof(int), cudaMemcpyDeviceToHost);
}
```

## Kernel definition

```
__global__ void AddIntegers(int *a, int *b,int *c)
{
    c[threadIdx.x]=a[threadIdx.x]+b[threadIdx.x];
}
```

**\_\_global\_\_** lets the compiler know that it is a kernel function and accordingly appropriate actions will be taken

Call to the kernel from the host is made using the below format:  
**<<<number of blocks, number of threads>>>**

## Built-in Variables

- In the kernel a set of built-in variables specifies the grid and block dimensions (Dim) and indices (Idx).
- These can be used to determine the thread ID
  - blockDim contains .x and .y grid dimensions (sizes)
  - blockIdx contains block indices .x and .y in the grid
  - blockDim contains the thread block .x, .y, .z dimensions (sizes) –  
threadIdx contains .x, .y and .z thread block indices
- 1D thread block: ID = threadIdx.x
- 2D thread block: ID = threadIdx.x + threadIdx.y \* blockDim.x
- 3D thread block: ID = threadIdx.x + threadIdx.y \* blockDim.x +  
threadIdx.z \* blockDim.x \* blockDim.y

## API functions in CUDA

- Device Memory Allocation function
- Host-Device Data /transfer

## Device Memory Allocation function

### cudaMalloc()

- Similar to Malloc function in C
- It will allocate memory in the GPU's global memory.
- Two parameters
  1. Address of a pointer to the allocated object
  2. Size of allocated object in terms of bytes

### cudaFree()

- Frees object from device global memory
- Parameter - Pointer to freed object

## Host-Device Data /transfer

### cudaMemcpy()

--Used for memory data transfer

Parameters:

1. Pointer to destination
2. Pointer to source
3. Number of bytes copied
4. Type of transfer

There are 4 types of data transfer:

-Host to Host

-Host to Device

-Device to Host

-Device to Device

All these data transfers are asynchronous.

## Programming Examples

### 2D Grid Block

- Cuda allows us to create 1-D,2-D,3-D grid blocks
- 2-D Kernel Launch:
- Image processing tasks typically impose a regular 2D raster (an image) over the problem domain. Computational Fluid dynamics tasks might be most naturally expressed by partitioning a volume over a 3D grid.

## 2-D Representation Program

- 2-D mapping

```
__global__ void kernel(int *array)
{
    int index_x = blockIdx.x * blockDim.x + threadIdx.x;
    int index_y = blockIdx.y * blockDim.y + threadIdx.y;

    // map the two 2D indices to a single linear, 1D index
    int grid_width = gridDim.x * blockDim.x;
    int index = index_x * grid_width + index_y;

    // map the two 2D block indices to a single linear, 1D block index
    int result = blockIdx.x * gridDim.x + blockIdx.y;

    // write out the result
    array[index] = result;
}
```

## Vector Addition Map

### 2:1 Mapping

#### Iteration Space

Range of variable i  
 Range of variable j  
 0....255  
 0....255

#### Data Space

Vectors A(2-D array), C  
 Function used f(i,j)  
 □□□□.....

#### Thread Space

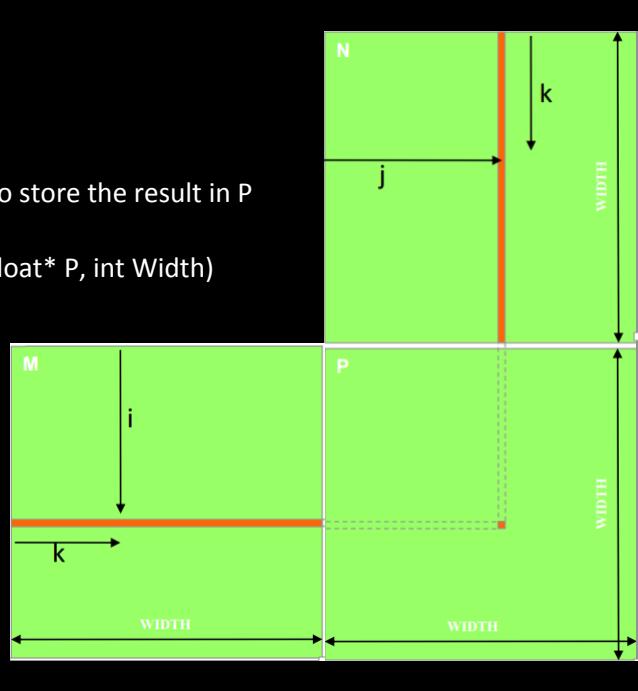
i and j will be mapped as  
 $f(i, \text{ThreadId}, \text{BlockId})$   
 △△△△ .....16  
 △△△△ ...16

There exists a 2:1 mapping between the iterator variable i, j and the ThreadId for every BlockId

There exists 16 blocks each being mapped to a set of 2 values in x and y direction

# Matrix Multiplication

- Consider two Matrices M and N. We need to store the result in P
- Executing on the host:
- `void MatrixMulOnHost(float* M, float* N, float* P, int Width)`
- {
- for (int i = 0; i < Width; ++i)
- for (int j = 0; j < Width; ++j) {
 double sum = 0;
 for (int k = 0; k < Width; ++k) {
 double a = M[i \* Width + k];
 double b = N[k \* Width + j];
 sum += a \* b;
 }
 P[i \* Width + j] = sum;
 }
- }



## Cuda Code:

```
void MatrixMultiplication(float* M, float* N, float* P, int Width)
{
    int size = Width * Width * sizeof(float);
    float* Md, Nd, Pd;

    1. // Transfer M and N to device memory
    cudaMalloc((void**)&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);

    // Allocate P on the device
    cudaMalloc((void**)&Pd, size);

    2. // Kernel invocation code - to be shown later
    ...
    3. // Transfer P from device to host
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

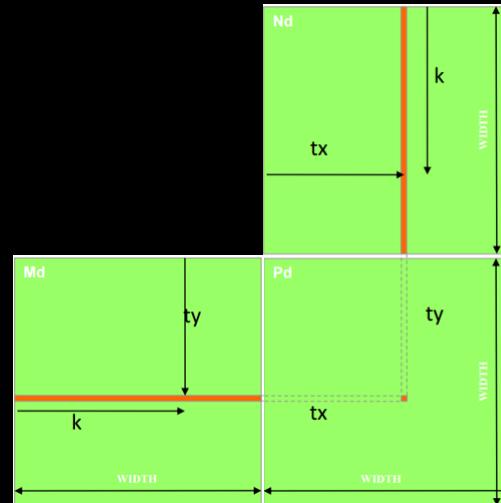
## Kernel Code

```
// Matrix multiplication kernel - thread specification
__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int W
{
    // 2D Thread ID
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    // Pvalue stores the Pd element that is computed by the thread
    float Pvalue = 0;

    for (int k = 0; k < Width; ++k)
    {
        float Mdelement = Md[ty * Width + k];
        float Ndelement = Nd[k * Width + tx];
        Pvalue += Mdelement * Ndelement;
    }

    // Write the matrix to device memory each thread writes one element
    Pd[ty * Width + tx] = Pvalue;
}
```



## Programming Example-Jacobi1D.cu

## References:

- <http://www.cc.gatech.edu/~vetter/keeneland/tutorial-2011-04-14/02-cuda-overview.pdf>
- <https://class.coursera.org>
- Programming Massively Parallel Processors - David B. Kirk and Wen-mei W. Hwu
- <http://docs.nvidia.com/cuda/parallel-thread-execution/#axzz3p0rtlrXV>
- <https://code.google.com/p/stanford-cs193g-sp2010/wiki/TutorialMultidimensionalKernelLaunch>