

Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.

# Parallel Programming in C with MPI and OpenMP

Michael J. Quinn



Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.

## Chapter 17

Shared-memory Programming

## Outline

- OpenMP
- Shared-memory model
- Parallel **for** loops
- Declaring private variables
- Critical sections
- Reductions
- Performance improvements
- More general data parallelism
- Functional parallelism

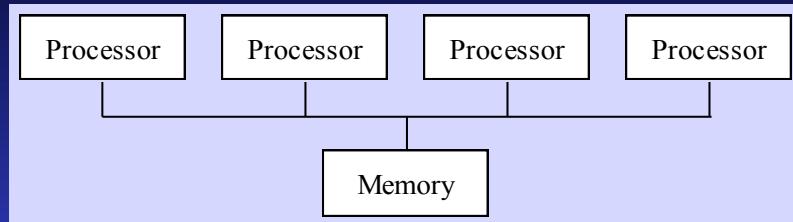
## OpenMP

- OpenMP: An application programming interface (API) for parallel programming on multiprocessors
  - ◆ Compiler directives
  - ◆ Library of support functions
- OpenMP works in conjunction with Fortran, C, or C++

## What's OMP/MPI/CUDA good for?

- C+OpenMP shared memory multiprocessors
- C+MPI(+OpenMP) distributed machines  
(built out of shared memory machines)
  - ◆ for us:  
ISTeC Cray:  
52 nodes, each with  
4 processors, each with 6 cores
- C+CUDA for GPUs
  - ◆ for us: Nvidia Tesla 30 processors, 240 cores, on a chip

## Shared-memory Model

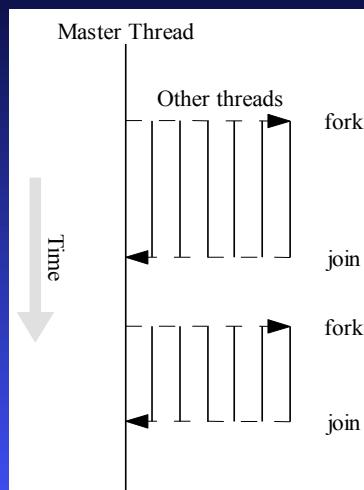


Processors interact and synchronize with each other through shared variables.

## Fork/Join Parallelism

- Initially only master thread is active
- Master thread executes sequential code
- Fork: Master thread creates or awakens additional threads to execute parallel code
- Join: At end of parallel code created threads die or are suspended

## Fork/Join Parallelism



## Incremental Parallelization

- Sequential program a special case of a shared-memory parallel program
- Parallel shared-memory programs may only have a single parallel loop
- Incremental parallelization: process of converting a sequential program to a parallel program a little bit at a time

## Parallel for Loops

- C programs often express data-parallel operations as **for** loops

```
for (i = first; i < size; i += prime)
    marked[i] = 1;
```
- OpenMP makes it easy to indicate when the iterations of a loop may execute in parallel
- Compiler takes care of generating code that forks/joins threads and allocates the iterations to threads

## Pragmas

- Pragma: a compiler directive in C or C++
- Stands for “pragmatic information”
- A way for the programmer to communicate with the compiler
- Compiler free to ignore pragmas
- Syntax:

**#pragma omp <rest of pragma>**

## Parallel for Pragma

- Format:  
**#pragma omp parallel for  
for (i = 0; i < N; i++)  
    a[i] = b[i] + c[i];**

- There is a data dependence between the iterations, but it can be resolved. HOW??

## Parallel for Pragma

- Format:

```
#pragma omp parallel for
for (i = 0; i < N; i++)
    a[i] = b[i] + c[i];
```

- There is a data dependence between the iterations, but it can be resolved. HOW??
  - ◆ Blocking: compute #iterations, group them , compute start index of each group.
- Compiler must be able to verify the run-time system will have information it needs to schedule loop iterations. What can we NOT have?

## Parallel for Pragma

- Format:

```
#pragma omp parallel for
for (i = 0; i < N; i++)
    a[i] = b[i] + c[i];
```

- There is a data dependence between the iterations, but it can be resolved. HOW??
  - ◆ Compute #iterations, group them , compute start index of each group.
- Compiler must be able to verify the run-time system will have information it needs to schedule loop iterations. What can we NOT have?
  - ◆ Dynamic control, e.g. conditional exit out of the loop.

## Canonical Shape of for Loop Control Clause

```
for(index = start; index  $\geq$  {<=  
           >= } end; {  
           index ++  
           + +index  
           index --  
           - -index  
           index+ = inc  
           index- = inc  
           index = index +inc  
           index = inc +index  
           index = index - inc  
           })}
```

## Execution Context

- Every thread has its own execution context
- Execution context: address space containing all of the variables a thread may access
- Contents of execution context:
  - ◆ static variables
  - ◆ dynamically allocated data structures in the heap
  - ◆ variables on the run-time stack
  - ◆ additional run-time stack for functions invoked by the thread

## Shared and Private Variables

- Shared variable: has same address in execution context of every thread
- Private variable: has different address in execution context of every thread
- A thread cannot access the private variables of another thread

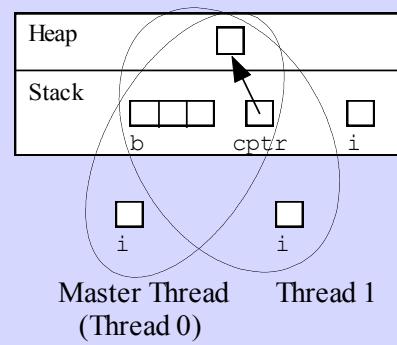
## C memory structure recap

- Data variables in C can be:
  - ◆ In registers (compiler directive, may not always be followed, and may spill)
  - ◆ On the stack (activation record of the current execution context)
  - ◆ On the heap – dynamically allocated through malloc

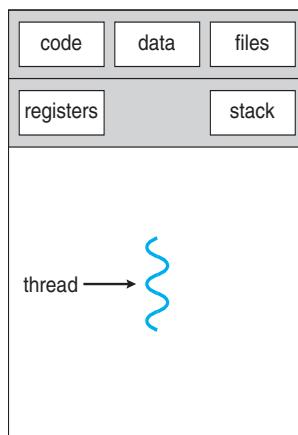
## Shared and Private Variables

```
int main (int argc, char *argv[])
{
    int b[3];
    char *cptr;
    int i;

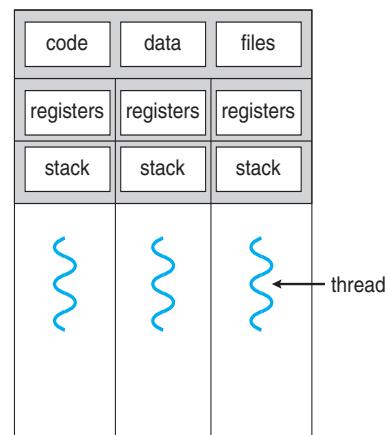
    cptr = malloc(1);
#pragma omp parallel for
    for (i = 0; i < 3; i++)
        b[i] = i;
```



## Single and Multithreaded Processes



single-threaded process



multithreaded process



## Function `omp_get_num_procs`

- Returns number of processors available for use by the parallel program

```
int omp_get_num_procs (void)
```

## Function `omp_set_num_threads`

- Uses the parameter value to set the number of threads to be active in parallel sections of code
- May be called at multiple points in a program

```
void omp_set_num_threads (int t)
```

- Can you set the #threads > OMP\_NUM\_THREADS?
  - ◆ try it

## Declaring Private Variables

```
for (i = 0; i < n; i++)
    for (j = 0; j < n; j++)
        a[i][j] = MIN(a[i][j], a[i][k]+tmp);
```

- Either loop could be executed in parallel
- We prefer to make outer loop parallel, to reduce number of forks/joins
- We then must give each thread its own private copy of variable **j**, **WHY?**

## private Clause

- Clause: an optional, additional component to a pragma
- Private clause: directs compiler to make one or more variables private

```
private ( <variable list> )
```

## Example Use of private Clause

```
#pragma omp parallel for private(j)
for (i = 0; i < n; i++)
    for (j = 0; j < n; j++)
        a[i][j] = MIN(a[i][j], a[i][k]+tmp);
```

## firstprivate Clause

- Used to create private variables having initial values identical to the variable controlled by the master thread as the loop is entered
- Variables are initialized once per thread, not once per loop iteration
- If a thread modifies a private variable's value in an iteration, subsequent iterations that are executed by **this thread** will get the modified value, other threads will work with their own private copy

## lastprivate Clause

- Sequentially last iteration: iteration that occurs last when the loop is executed sequentially
- **lastprivate** clause: used to copy back to the master thread's copy of a variable the private copy of the variable from the thread that executed the **sequentially** last iteration

## Race Condition

- Consider this C program segment to compute  $\pi$  using the rectangle rule:

```
double area, pi, x;
int i, n;
...
area = 0.0;
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
    area += 4.0/(1.0 - x*x);
}
pi = area / n;
```

## Race Condition (cont.)

- If we simply parallelize the loop...

**what happens?**

```
double area, pi, x;  
int i, n;  
...  
area = 0.0;  
#pragma omp parallel for private(x)  
for (i = 0; i < n; i++) {  
    x = (i+0.5)/n;  
    area += 4.0/(1.0 - x*x);  
}  
pi = area / n;
```

## Race Condition (cont.)

- ... we set up a race condition in which two processes may both read **area**, then compute the rhs, then write back

**area**

15.230

**Answer should be 18.995**

Thread A

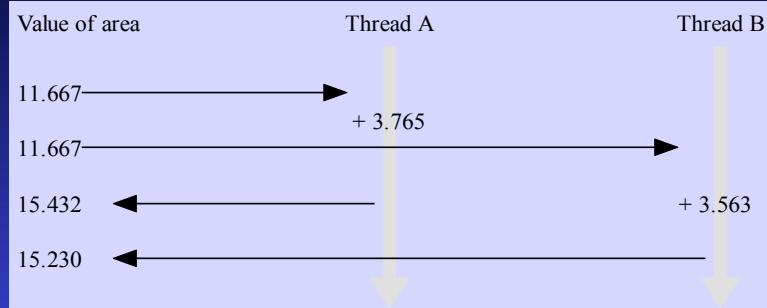
15.432

Thread B

15.230

**area += 4.0 / (1.0 - x\*x)**

## Race Condition Time Line



## critical Pragma

- Critical section: a portion of code that only one thread at a time may execute
- We denote a critical section by putting the pragma

**#pragma omp critical**

in front of a block of C code

## Correct, but inefficient, code

```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x)
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
#pragma omp critical
    area += 4.0/(1.0 - x*x);
}
pi = area / n;
```

## Source of Inefficiency

- Update to **area** inside a critical section
- Only one thread at a time may execute the statement; i.e., it is sequential code
- Time to execute statement significant part of loop, made worse by the need for synchronization
- Speedup will be severely constrained
- Can you see a solution?

## Reductions

- Reductions are so common that OpenMP provides support for them
- May add reduction clause to **parallel for** pragma
- Specify reduction operation and reduction variable
- OpenMP takes care of storing partial results in private variables and combining partial results after the loop

## reduction Clause

- The reduction clause has this syntax:  
**reduction (<op> :<variable>)**
- Operators
  - ◆ + Sum
  - ◆ \* Product
  - ◆ & Bitwise and
  - ◆ | Bitwise or
  - ◆ ^ Bitwise exclusive or
  - ◆ && Logical and
  - ◆ || Logical or

## π code with Reduction Clause

```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for \
    private(x) reduction(+:area)
for (i = 0; i < n; i++) {
    x = (i + 0.5)/n;
    area += 4.0/(1.0 - x*x);
}
pi = area / n;
```

## Performance Improvement #1

- Too many fork/joins can lower performance
- Inverting loops may help performance if
  - ◆ Parallelism is in inner loop
  - ◆ After inversion, the outer loop can be made parallel
  - ◆ Inversion does not significantly lower cache hit rate

## Performance Improvement #2

- If loop has too few iterations, fork/join overhead is greater than time savings from parallel execution
- The **if** clause instructs compiler to insert code that determines at run-time whether loop should be executed in parallel; e.g.,

```
#pragma omp parallel for if(n > 5000)
```

## Performance Improvement #3

- We can use **schedule** clause to specify how iterations of a loop should be allocated to threads
- Static schedule: all iterations allocated to threads before any iterations executed
- Dynamic schedule: only some iterations allocated to threads at beginning of loop's execution. Remaining iterations allocated to threads that complete their assigned iterations.

## Static vs. Dynamic Scheduling

- Static scheduling
  - ◆ Low overhead
  - ◆ May exhibit high workload imbalance
- Dynamic scheduling
  - ◆ Higher overhead
  - ◆ Can reduce workload imbalance

## Chunks

- A chunk is a contiguous range of iterations
- Increasing chunk size reduces overhead and may increase cache hit rate
- Decreasing chunk size allows finer balancing of workloads

## schedule Clause

- Syntax of schedule clause  
**schedule (<type>[ ,<chunk> ])**
- Schedule type required, chunk size optional
- Allowable schedule types
  - ◆ static: static allocation
  - ◆ dynamic: dynamic allocation
  - ◆ guided
  - ◆ runtime: type chosen at run-time based on value of environment variable OMP\_SCHEDULE

## Scheduling Options

- schedule(static): block allocation of about  $n/\#threads$  contiguous iterations to a thread
- schedule(static,C): interleaved allocation of chunks of size C to threads
- schedule(dynamic): dynamic one-at-a-time allocation of iterations to threads
- schedule(dynamic,C): dynamic allocation of C iterations at a time to threads

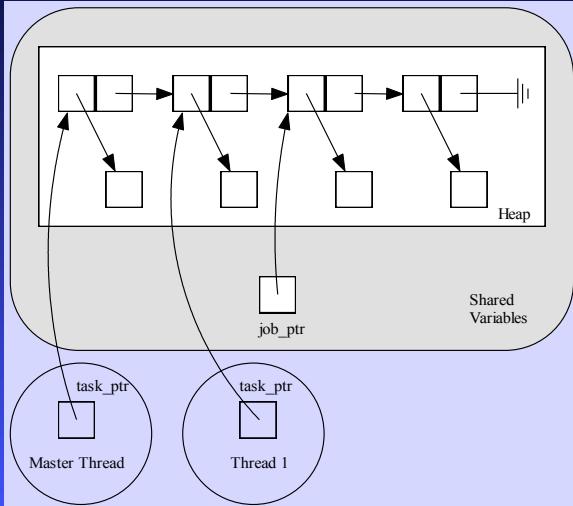
## Scheduling Options (cont.)

- `schedule(guided, C)`: dynamic allocation of chunks to tasks using guided self-scheduling heuristic. Initial chunks are bigger, later chunks are smaller, minimum chunk size is  $C$ .
- `schedule(guided)`: guided self-scheduling with minimum chunk size 1
- `schedule(runtime)`: schedule chosen at run-time based on value of `OMP_SCHEDULE`; Unix example:  
`setenv OMP_SCHEDULE "static,1"`

## More General Data Parallelism

- Our focus has been on the parallelization of **for** loops
- Other opportunities for data parallelism
  - ◆ processing items on a “to do” list
  - ◆ **for** loop + additional code outside of loop

## Processing a “To Do” List



## parallel Pragma

- The **parallel** pragma precedes a block of code that should be executed by *all* of the threads
- Note: execution is replicated among all threads

## Use of **parallel** Pragma

```
#pragma omp parallel private(task_ptr)
{
    task_ptr = get_next_task (&job_ptr);
    while (task_ptr != NULL) {
        complete_task (task_ptr);
        task_ptr = get_next_task (&job_ptr);
    }
}
```

## Critical Section for **get\_next\_task**

```
char *get_next_task(struct job_struct
                    **head) {
    struct task_struct *nextTask;
    #pragma omp critical
    {
        if (*head == NULL) nextTask = NULL;
        else {
            nextTask = (*head)->task;
            *head = (*head)->next;
        }
    }
    return nextTask;
}
```

## Functions for SPMD-style Programming

- The parallel pragma allows us to write SPMD-style programs
- In these programs we often need to know number of threads and thread ID number
- OpenMP provides functions to retrieve this information

### Function `omp_get_thread_num`

- This function returns the thread identification number
- If there are  $t$  threads, the ID numbers range from 0 to  $t-1$
- The master thread has ID number 0

```
int omp_get_thread_num (void)
```

## Function omp\_get\_num\_threads

- Function `omp_get_num_threads` returns the number of active threads
- If call this function from sequential portion of program, it will return 1

```
int omp_get_num_threads (void)
```

## for Pragma

- The **parallel** pragma instructs every thread to execute all of the code inside the block
- If we encounter a **for** loop that we want to divide among threads, we use the **for** pragma

```
#pragma omp for
```

## Example Use of for Pragma

```
#pragma omp parallel private(i,j)
for (i = 0; i < m; i++) {
    low = a[i];
    high = b[i];
    if (low > high) {
        printf ("Exiting (%d)\n", i);
        break;
    }
#pragma omp for
    for (j = low; j < high; j++)
        c[j] = (c[j] - a[i])/b[i];
}
```

## single Pragma

- Suppose we only want to see the output once
- The **single** pragma directs compiler that only a single thread should execute the block of code the pragma precedes
- Syntax:

```
#pragma omp single
```

## Use of single Pragma

```
#pragma omp parallel private(i,j)
for (i = 0; i < m; i++) {
    low = a[i];
    high = b[i];
    if (low > high) {
#pragma omp single
        printf ("Exiting (%d)\n", i);
        break;
    }
#pragma omp for
    for (j = low; j < high; j++)
        c[j] = (c[j] - a[i])/b[i];
}
```

## nowait Clause

- Compiler puts a barrier synchronization at end of every parallel for statement
- In our example, this is necessary: if a thread leaves loop and changes **low** or **high**, it may affect behavior of another thread
- If we make these private variables, then it would be okay to let threads move ahead, which could reduce execution time

## Use of nowait Clause

```
#pragma omp parallel private(i,j,low,high)
for (i = 0; i < m; i++) {
    low = a[i];
    high = b[i];
    if (low > high) {
#pragma omp single
        printf ("Exiting (%d)\n", i);
        break;
    }
#pragma omp for nowait
    for (j = low; j < high; j++)
        c[j] = (c[j] - a[i])/b[i];
}
```

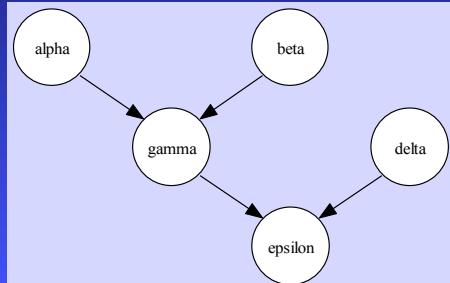
## Functional Parallelism

- To this point all of our focus has been on exploiting data parallelism
- OpenMP allows us to assign different threads to different portions of code (functional parallelism)

## Functional Parallelism Example

```
v = alpha();  
w = beta();  
x = gamma(v, w);  
y = delta();  
printf ("%6.2f\n", epsilon(x,y));
```

May execute alpha,  
beta, and delta in  
parallel



## parallel sections Pragma

- Precedes a block of  $k$  (sub) blocks of code that may be executed concurrently by (up to)  $k$  threads
- Syntax:

```
#pragma omp parallel sections
```

## section Pragma

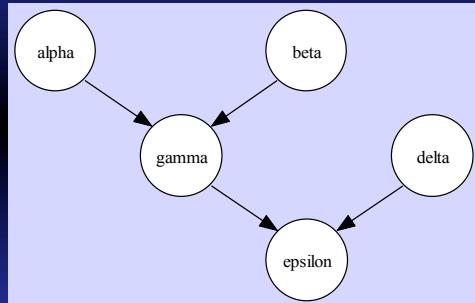
- Precedes each block of code within the encompassing block preceded by the parallel sections pragma
- May be omitted for first parallel section after the parallel sections pragma
- Syntax:

```
#pragma omp section
```

## Example of parallel sections

```
#pragma omp parallel sections
{
    #pragma omp section /* Optional */
        v = alpha();
    #pragma omp section
        w = beta();
    #pragma omp section
        y = delta();
}
x = gamma(v, w);
printf ("%6.2f\n", epsilon(x,y));
```

## Another Approach



Execute alpha and  
beta in parallel.  
Execute gamma and  
delta in parallel.

## sections Pragma

- Appears inside a parallel block of code
- Has same meaning as the **parallel sections** pragma
- If multiple **sections** pragmas inside one parallel block, may reduce fork/join costs

## Use of sections Pragma

```
#pragma omp parallel
{
    #pragma omp sections
    {
        v = alpha();
        #pragma omp section
        w = beta();
    }
    #pragma omp sections
    {
        x = gamma(v, w);
        #pragma omp section
        y = delta();
    }
}
printf ("%6.2f\n", epsilon(x,y));
```

## Summary (1/3)

- OpenMP is an API for shared-memory parallel programming
- Shared-memory model based on fork/join parallelism
- Data parallelism
  - ◆ parallel for pragma
  - ◆ reduction clause

## Summary (2/3)

- Functional parallelism (parallel sections pragma)
- SPMD-style programming (parallel pragma)
- Critical sections (critical pragma)
- Enhancing performance of parallel for loops
  - ◆ Inverting loops
  - ◆ Conditionally parallelizing loops
  - ◆ Changing loop scheduling

## Summary (3/3)

<i>Characteristic</i>	<i>OpenMP</i>	<i>MPI</i>
Suitable for multiprocessors	Yes	Yes
Suitable for multiccomputers	No	Yes
Supports incremental parallelization	Yes	No
Minimal extra code	Yes	No
Explicit control of memory hierarchy	No	Yes